

Marianne Arriola

New York, NY | marriola@cs.cornell.edu | m-arriola.com

Interests: Deep Learning, Generative AI for text, Discrete Diffusion

Education

Ph.D. student in Computer Science, Cornell University Sep 2023–Jun 2028

Advisor: Volodymyr Kuleshov

B.S., Computer Science, University of California, Santa Barbara Aug 2019–Jun 2023

Advisor: Ambuj Singh

Current Research

Diffusion Language Models

Committee: Volodymyr Kuleshov (Chair), Mohamed Abdelfattah, Fei Wang

- I explore diffusion LMs which enable faster & more controllable generation with better benchmark performance for math & planning compared to traditional LLMs
- I design novel parameterizations and architectures to improve diffusion LM quality [1,2], training efficiency [1], and inference capabilities [1]

Publications

- [1] [Marianne Arriola](#), Aaron Gokalsan, Justin Chiu, Zhihan Yang, Zhixuan Qi, Jiaqi Han, Subham Sahoo, Volodymyr Kuleshov. "[Block Diffusion: Interpolating between Autoregressive and Diffusion Language Models](#)." *ICLR 2025*. **Oral presentation (Top 1.77%)**.
- [2] Subham Sahoo, [Marianne Arriola](#), Yair Schiff, Aaron Gokaslan, Edgar Marroquin, Justin Chiu, Alexander Rush, Volodymyr Kuleshov "[Simple and Effective Masked Diffusion Language Models](#)." *NeurIPS 2024*.
- [3] [Marianne Arriola](#), Weishen Pan, Manqi Zhou, Qiannan Zhang, Chang Su, Fei Wang. "[Joint Analysis of Single-Cell Data across Cohorts with Missing Modalities](#)." *arXiv* (Feb 2024).
- [4] [Marianne Arriola](#) & Kadina Johnston "[Identifying Optimal Proteins by Their Structure Using Graph Neural Networks](#)." *Caltech URJ* (Jun 2022).

Employment

Runway, Research Intern. New York, NY Jun 2025–Aug 2025

MIT CSAIL, Research Intern. Cambridge, MA Jun 2022–Nov 2022

Caltech, Research Intern. Caltech, Pasadena, CA Jun 2021–Aug 2021

Projects

MIT CSAIL, Cambridge, MA Jun 2022 – Nov 2022

PI: Justin Solomon

- Designed a memory-efficient representation of complex point clouds (i.e. from LiDAR) that summarizes points using geometric primitives

- Developed a graph neural network for hybrid point clouds that achieves comparable segmentation performance to state-of-the-art methods while halving memory requirements

Caltech, Pasadena, CA
PI: Frances Arnold

Jun 2021 – Aug 2021

- Proposed a data-driven method to iteratively refine existing protein structures for desired properties (i.e. substrate specificity)
- Built a graph neural network to predict protein functional capacity using graph-based structure representations

Open-Source Contributions

Block Diffusion Language Models ([GitHub](#))

- Led development of a diffusion LM with arbitrary-length generation and KV caching
- 700 stars as of Jul 2025

Patents

Marianne Arriola*, Xiangru Huang*, Yue Wang, Vitor Campagnolo Guizilini, Rares Andrei Ambrus, Justin Solomon. "Hybrid Geometric Primitive Representations for Point Clouds." U.S. patent pending.

Selected Talks

International Conference on Learning Representations (ICLR)	Apr 2025
Amazon Artificial General Intelligence (AGI)	Apr 2025

Awards

NSF Graduate Research Fellowship	Mar 2023–Jun 2028
Bowers CIS Dean's Excellence Fellowship, Cornell University	Mar 2023–Jun 2029

Activities

MIT Summer Research Program Application Reviewer	Jan 2024–Feb 2024
--	-------------------

Programming Skills

Python, PyTorch, C++, MATLAB, Bash